



**Politecnico
di Torino**

Operating Systems for Embedded Systems

Project's material overview

DMA Attacks - Group 4

Elias Michael - s307276

Marino Lorenzo - s317703

Sansone Giacomo - s307761

Valenza Arianna - s317742

Contents

1	Structure of the work	2
2	Overview on DMA attacks	3
2.1	Prerequisites	3
2.1.1	What is the DMA	3
2.1.2	Memory protection	4
2.2	An introduction to DMA attacks	4
2.3	Some examples of DMA attacks	5
2.4	Countermeasures	5
3	Documents from literature	7
3.1	Paper 1	7
3.2	Paper 2	9
3.3	Paper 3	10
3.4	Paper 4	11
4	Documentation	13
4.1	UART Library	13
4.1.1	uart_parameters	13
4.1.2	init_uart	15
4.2	DMA Library	15
4.2.1	dma_transfer	15
4.2.2	dma_lli_element	16
4.2.3	init_DMA	16
4.2.4	transfer_DMA	16
4.3	Application	17
4.3.1	Setup	17
4.3.2	Micrium OS	17
4.3.3	APP task	17
4.3.4	Lock task	18
4.3.5	Display task	19
4.3.6	DMA_IRQHandler	19

Chapter 1

Structure of the work

This document aims to be a collection of all the material we got and developed for the project of *Operating Systems for Embedded Systems*.

The request was to collect the material in a way that is easily accessible by a student who is studying this topic. From our point of view, the most useful thing to do was to mix some explanations with other external resources. In this way, we can ensure to provide the correct information before introducing the material we have used.

This document has the following structure:

- chapter 2 is a simple introduction to Direct Access Memory (DMA) attacks. This type of attack is not particularly difficult to understand, as the interesting part is to perform (and try to avoid) it, so that is only a small part of the work;
- chapter 4 explains the two libraries we wrote for the LPC1768 microprocessor, to activate and use the UART peripheral and the GPDMA controller;
- chapter 3 presents brief overviews of some of the relevant papers that we have identified on this topic.

This work is licensed under a Creative Commons “Attribution-NonCommercial-ShareAlike 4.0 International” license.



Chapter 2

Overview on DMA attacks

2.1 Prerequisites

2.1.1 What is the DMA

Direct Memory Access is a feature that allows different peripheral, to access the memory directly, independently of the CPU. This frees up the CPU to perform other tasks and can result in a significant performance boost for the system. Many hardware systems use DMA, including disk drive controllers, graphics cards and network cards. It is an hardware technology, so it works independently from the operating system, and therefore, since it is working with physical addresses, it represent a source of non deterministic behaviors, and a possible backdoor for malware attacks. DMA is typically composed of:

- A DMA controller: responsible for coordinating the transfer of data between memory and peripheral devices.
- A set of DMA channels: allows multiple devices to perform DMA transfers simultaneously.
- A DMA buffer: a small amount of memory reserved for use by the DMA controller to temporarily store data during transfers.

We can have two different organization of the DMA engine:

- Third-party DMA: The host CPU configures (source and destination address) the central DMA controller to perform a DMA transfer. The DMA controller interrupts the host CPU when the DMA transfer has been finished. Hence, the host CPU is aware of a third-party DMA transfer.
- First-party DMA: The peripheral device can configure its own DMA engine. The device acts as bus master to get control of the system

bus to perform a DMA transfer. The device can interrupt the host CPU when the device has completed the transfer. The transfer also works if the device does not interrupt the host CPU at the end of the DMA transfer. In this case the CPU is completely unaware of the DMA transfer.

2.1.2 Memory protection

Memory protection is a crucial aspect of computer systems, especially for modern operating systems, as it helps prevent errors and security breaches. Instead of using the real addresses of the memory space, we use *virtual addresses* which will be translated to the real ones by proper hardware devices. Since there is always an entity (called MMU - Memory Management Unit) between us and the memory, it is impossible to access freely the memory. It can be implemented in different ways, such as:

- Memory segmentation: This technique divides memory into distinct segments, each with its own specific purpose and access control permissions. This helps prevent programs from accessing memory locations they are not authorized to access.
- Paging: Paging is a technique where the physical memory is divided into fixed-sized blocks called pages, which can be individually assigned to different processes. This allows the operating system to control and monitor the memory access of each process.
- Access control: Access control is a mechanism that defines which processes or users are allowed to access specific memory locations, and what actions they are permitted to perform. This helps prevent unauthorized access to sensitive information.
- Virtual memory: Virtual memory is a technique that allows a process to use more memory than physically available by temporarily transferring data from physical memory to a disk. This helps ensure that processes don't interfere with each other's memory usage.

The implementation can vary depending on the operating system, hardware, and security requirements.

In spite of its many advantages, in many embedded systems this implementation is avoided to reduce costs, complexity and time to market.

2.2 An introduction to DMA attacks

DMA attacks are a particularly powerful class of attacks, especially for embedded systems. These attacks enable a potential attacker to read

and write memory off a victim system directly, bypassing the main CPU and OS. Achieving DMA attacks in embedded systems is easier than in a general purpose computer, due to the absence of memory protection, and more dangerous because often flat architecture operating systems are implemented in those kind of systems, which allow the attacker to access reserved and critical areas, like the ones of the OS itself.

By overwriting memory in a DMA transfer, for instance just increasing the size of the burst transfer, the attacker can modify and control reserved memory regions, performing any manner of malicious activity.

Direct Memory Access is a capability implemented in every system which has the need of speeding up the transfer from/to peripherals. It permits to execute the transfer without the aid of neither CPU nor Operating System, after setting up the DMA, so that the peripherals can directly communicate with the system's memory and save time.

Despite the possibility to accelerate the system, DMA exposes the system to various attacks, for instance:

- Access and/or modify reserved information;
- Extend control over the execution of the kernel itself;
- Malware injection.

2.3 Some examples of DMA attacks

One of the main tools used to perform DMA attacks is pcileech. Most of the works you can find on the internet regarding DMA attacks are proof of concept, though there are some cases in which hackers were able to inject malicious code inside a computer once they had access to the hardware.

2.4 Countermeasures

DMA attacks can be difficult to prevent because they occur at a hardware level and are not visible to the operating system or CPU. However, there are some measures that can be taken to reduce the risk of DMA attacks and make them more difficult to execute.

One way to prevent DMA attacks is to use an IOMMU (Input/Output Memory Management Unit). An IOMMU is a hardware component that acts as a memory management unit for input/output (I/O) operations. It allows the system to restrict the memory addresses that I/O devices can access, making it more difficult for an attacker to perform a DMA attack.

A different approach is to use hardware-enforced memory isolation. This technique involves dividing the system's memory into different regions, with each region assigned a specific level of access. This can help to prevent an attacker from gaining access to sensitive areas of memory.

Another measure that can be considered to prevent DMA attacks is to use secure boot. Secure boot is a process that verifies the integrity of the system's boot process. It ensures that only trusted software is allowed to run when the system boots up, making it more difficult for an attacker to inject malicious code into the system.

Other approaches to preventing DMA attacks include using hardware-based firewalls, implementing role-based access control, and using hardware-based encryption.

Chapter 3

Documents from literature

We used the following papers understand better the way DMA attacks are performed.

3.1 Paper 1

The title of the paper is *Characterizing, exploiting, and detecting DMA code injection vulnerabilities in the presence of an IOMMU* [1].

This paper primarily focuses on the challenges of conducting a DMA attack when an Input/Output (IO) Memory Management Unit (MMU) is present (an IO MMU is similar to a regular MMU but only deals with IO accesses). If you are not familiar with virtual memory on x86 systems, this paper may be a bit difficult to understand. Since our system does not have virtual memory or protection, our goal is to summarize some of the key ideas from the paper using only physical address space.

The first half of the paper is the most important, as it provides an overview of DMA attacks and the necessary conditions to perform a code injection. In the second half, the authors present a tool they developed to perform static analysis of the Linux kernel and identify vulnerabilities of a particular type.

A system with a DMA controller is always vulnerable to DMA attacks. These attacks can occur concurrently with normal CPU execution, making it difficult for the system to detect them. DMA attacks can often be executed by exploiting firmware bugs. There are various types of attacks that can be carried out using DMA, including injecting and executing code on the target system, stealing sensitive information such as secret cipher keys, and dumping the entire memory.

As previously mentioned, the paper discusses some characteristics of modern operating systems that can be helpful in defending against DMA attacks. We will briefly summarize them here to explain why they are not relevant in our context.

- IOMMU: an Input/Output Memory Management Unit (IOMMU) allows peripherals to access only certain virtual addresses (IOVAs) and prevents them from accessing physical addresses directly. A single physical address can be mapped to multiple IOVAs, just as in the CPU's virtual memory. In our system, we don't use virtual addresses, as we do not have a MMU: both the CPU and the peripherals always use physical addresses without any form of protection.
- NX-BIT: in a system with virtual memory, pages have certain bits that describe the operations that can be performed on them. For example, if the W bit is 1, you can write to that page using a certain virtual address (VA). If you try to perform this operation on a page with the W bit set to 0, you will receive an exception (likely a segmentation fault). One bit that is present is the NX bit, which determines whether the content of that page can be used as executable code or not. If a peripheral is able to inject malicious code into memory and that area of memory has NX=0, the code cannot be executed. If you try to execute it, you will receive an exception. This is not an issue in our system, since we do not have any protection. If we can write to the Program Counter register, we can execute malicious code from any location in memory as long as we have write access.
- (K)ASLR: in many DMA attacks, we need to know a specific memory address in order to steal data or inject malicious code. This is easy if we already know the memory layout. Some systems randomize the memory layout and the location of important data structures at each boot to prevent attackers from guessing the locations based on previous execution. Since our system does not have this feature, we can assume that it may be easy to locate addresses by using the .mem file after compiling/flashing the board.

The authors of the paper often refer to sub-page vulnerabilities, which occur when the target buffer of a DMA transaction is smaller than a page and other data on that page may be modified or stolen. There is a classification of attacks based on this characteristic. Since we are not concerned with pages, we can think about this classification in relation to physical memory only:

- The IO buffer (target of the DMA transaction) is part of a larger data structure that also includes function pointers. These pointers may be used as function callbacks to execute our malicious code.
- There are some metadata of the operating system in addition to the IO buffer.

- The IO buffer is next to some dynamically allocated data structure. This is difficult to discover since we are not responsible for managing the heap.

To perform a code injection, we need the following three attributes:

1. A buffer filled with malicious code. If we are able to write to that buffer, we also know its physical address.
2. Write access to a function callback pointer, which can alter the CPU control flow and cause it to execute the malicious code.
3. A time window exists during which the device can modify the callback pointer and the CPU will subsequently jump to the pointed code.

3.2 Paper 2

The title of the paper is *Stealthy Information Leakage Through Peripheral Exploitation in Modern Embedded Systems* [4].

One thing that the authors emphasize at the beginning of the work is that security is often overlooked in the embedded domain because time-to-market is considered more valuable (and engineers must focus on implementing functionalities rather than making them secure). One way to detect infiltrations in an embedded device is to profile the CPU: if its usage is higher than usual, there may be malware inside. This is why DMA attacks are powerful: DMA does not interact with the CPU, so its activity is hidden from any profiling tools within the system (the suggested solution is to detect malware by examining DMA transactions).

To make the attack as general as possible, the authors use the Device Tree, a data structure provided by Embedded Linux that contains information about all connected devices. By parsing its content, they can automate the process of choosing the appropriate output analog device and initialize the DMA transactions with the correct values.

To carry out the attack, the firmware of the board must be modified, which can be done physically (by changing the content of the boot memory) or remotely (by accessing the board through a network connection and gaining root access). There are various ways to verify the integrity of the original code, such as checking the hash of the entire code against a value stored in a secure memory (such as ROM, which cannot be modified). x86 systems are introducing the Secure Boot system, but it cannot be implemented in already operational embedded systems. This system is unlikely to be used in all new devices due to the area and power constraints in this domain.

After introducing the new kernel with our malware (some code that is periodically executed to exfiltrate data), we are almost finished. The attacker must decide which memory addresses to leak after conducting reconnaissance and considering the attack objectives. In general, embedded systems do not implement Address Space Layout Randomization (ASLR), so it is easy to find application data addresses such as the control runtime process and leak control-related information.

The remaining sections of the article discuss the actual implementation, mitigation techniques, and experimental results that demonstrate that DMA attacks do not affect CPU time.

3.3 Paper 3

The title of the paper is *Attacking TrustZone on devices lacking memory protection* [2].

ARM Trustzone offers a Trusted Execution Environment (TEE) embedded into the processor core. It establishes a new CPU operating mode called "secure mode", therefore it creates a distinction between a "secure" world and a "normal" world. The separation is artificial: both environments use the same CPU and RAM, the distinction is given by the Non-Secure bit (NS).

In this paper the authors exploit missing hardware feature on real embedded systems to execute arbitrary code in secure world: in order to work correctly TrustZone needs additional dedicated hardware controller, i.e. TZASC (TrustZone Address Space Controller) and TZPC (TrustZone Protection Controller).

The attack is performed on a Raspberry Pi 3 Model B. This particular board has a SoC with ARM11 core and TrustZone support, but lacks of the controller mentioned above.

On this SoC the authors run OP-TEE (a Trusted Execution Environment that implements ARM TrustZone) without NS-bit support, therefore the secure memory is accessible through DMA transactions. Their goal is to change the opcodes that return error values of key functions without changing the stack, in this way they can bypass OP-TEE OS Trusted Application (TA) signature validation and gain control of every TA of the system. To reach this goal, first they read the memory pages from the RAM, then compare this memory with compiled version of the secure OS and ARM Trusted Firmware to locate similar functions. Because the OS uses Address Space Layout Randomization, they were able to use the address from the memory dump to override trusted applications signature validations through the DMA, after doing so the TEE OS succeeded to validate their malicious TA, making possible the arbitrary code execution.

3.4 Paper 4

The provided document is a chapter of a conference proceeding of the 9th *GI International Conference on Detection of Intrusions and Malware & Vulnerability Assessment*, the chapter name is *Understanding DMA Malware*[3].

In this paper the authors explain how they implemented a malware exploiting DMA on x86 architectures, in particular executing malicious code on the processor provided by *Intel's Manageability Engine*.

The malware is called **DAGGER** (**Dm**A based keystroke **loGGER**) and in the implementations described in the paper it has been used to find and monitor the keyboard buffer.

In order to perform this kind of attack, controlling the DMA engine is not enough, because the target memory address usually is not known. The attacker must overcome *address randomization*, since the data structures are being saved in different location at every reboot of the system. Moreover the OS works with virtual addresses whether the DMA uses physical ones: the memory mapping is saved in so called *page tables*, which physical address is saved in a register not visible to the DMA engine, therefore the attacker needs to find a way to restrict the search space, otherwise it need to scan the whole host memory. In the document the authors suggest to analyze if the OS places the desired data structures in approximately the same region or to implement OS memory management mechanisms.

DAGGER performs a stealthy attack structured in three stages corresponding to three components:

1. Search: find the address of valuable data in the host memory via DMA.
2. Process Data: read valuable data within the regions identified in the search process.
3. Exfiltration: exfiltrate information in a way that is invisible to the host (e.g. using the Network Interface Card).

The authors implemented two prototypes to attack two different OS target: Windows and Linux. The search phase differs according to the OS to attack:

- Linux: the search phase is based on a signature scan which is compared with a signature derived analyzing the source code of the OS. The scan is restricted only on the kernel space, which is found in the first gigabyte of the system RAM. Exploiting GRUB, the kernel identifier is found at a constant physical memory address, using

this value it is possible to derive the offset used by the OS to map the virtual addresses to the physical ones. The authors successively restrict further the search space by empirically analyzing in which memory area the kernel places the needed data structure, which in this case is *USB Request Block*, in particular the field *transfer_dma*.

- Windows: differently from Linux, the memory mapping does not just use an offset, but *page tables* created by the kernel. The authors determined the physical address of this tables empirically and traverse them looking for a structure named *DeviceExtension* containing the last-pressed key buffer. The algorithm to find a starting point for the search it is more elaborate with respect to the previous case and requires more steps.

In the paper, in addition to the performance results of the malware, some possible countermeasure are discussed. It is claimed that the attack has not been detected by any anti-virus or anti-malware software present on the OS. As for the I/OMMU, the researchers enabled Intel VT-d and successfully blocked the attack on Linux, but not on Windows. They also state that to prevent the attack an I/OMMU needs an appropriate configuration and additional hardware that prevents the same malware from changing its configuration.

Chapter 4

Documentation

In this chapter the developed project simulating the DMA Attack through UART peripheral will be described in details. All the code developed is intended to be used with the microprocessor LPC1768. The user manual for this microprocessor can be found here. This was the main source for the project: as it is known, firmware development mainly consists in understanding and applying concepts from it, with the aim of a working system. Other interesting sources to write our code were this and this (though part of this presentation is in spanish).

The code is meant to be clear and modular, through the use of macros.

4.1 UART Library

The following subsections describe the composition of the UART library and how to enable the peripheral for communication (both sending and receiving). This library can be used with peripherals 0, 2, and 3. UART 1 has a different structure.

4.1.1 `uart_parameters`

```
1 struct uart_parameters{
2     uint32_t    uart_n;
3     uint32_t    clock_selection;
4     uint32_t    DLL;
5     uint32_t    DLM;
6     uint32_t    DivAddVal;
7     uint32_t    MulVal;
8     uint32_t    fifo_enable;
9     uint32_t    dma_mode;
10    uint32_t    trigger_level;
11    uint32_t    rdv_int;
12    uint32_t    word_size;
13    uint32_t    stop_bit_n;
14    uint32_t    parity_active;
15    uint32_t    parity_type;
16 };
```

Listing 4.1: Structure definition for UART parameters

This structure is used to define the communication parameters. Each field of the structure can take either a numeric value or a specific value as specified in the LPC1768 user manual. To make it easier for the user to use the library without needing to know all the low-level details, we provide many macros with the different values.

- **uart_n**: this value selects which of the three available UART peripheral you want to use. Possible values are `UART0`, `UART1` and `UART3`.
- **clock_selection**: these values determine which clock frequency the peripheral uses, selected from four different submultiples of the CPU clock (CCLK): `UART_CCLK_DIV_4` ($CCLK / 4$), `UART_CCLK_DIV_1` (CCLK), `UART_CCLK_DIV_2` ($CCLK / 2$), and `UART_CCLK_DIV_8` ($CCLK / 8$).
- **DLL, DLM, DivAddVal, MulVal**: these four values are used to set the baud rate for the UART connection, as explained in the LPC1768 user manual.
- **fifo_enable**: this value potentially enables the FIFO mode, which is required to use DMA. Its values are `UART_EN_FIFO` and `UART_NO_FIFO`.
- **dma_mode**: this value potentially enables the DMA mode. Its values are `UART_DMA` and `UART_NO_DMA`.
- **trigger_level**: the trigger level is the number of elements that need to be received before an interrupt is raised and in case the DMA is enabled, it will trigger a DMA request interrupt, that will result into a DMA transfer. Possible values are `UART_TR_LEVEL_1` (trigger level 1), `UART_TR_LEVEL_4` (trigger level 4), `UART_TR_LEVEL_8` (trigger level 8), and `UART_TR_LEVEL_16` (trigger level 16).
- **rdv_int**: with this value, we can enable an interrupt when new data is received. Possible values are `UART_RECV_DV_INT` and `UART_NO_INT`;
- **word_size**: this value determines the size of a word in the communication; possible values are `UART_WORD_5` (5), `UART_WORD_6` (6), `UART_WORD_7` (7) and `UART_WORD_8` (8).
- **stop_bit_n**: this value determines the number of stop bits we want, between 1 (`UART_STOP_1`) and 2 (`UART_STOP_2`).
- **parity_active**: this element allows to enable (`UART_PARITY`) or disable (`UART_NO_PARITY`) the parity field in the transmission.
- **parity_type**: this determines the type of parity we want, between odd (`UART_ODD_PARITY`) and even (`UART_EVEN_PARITY`).

4.1.2 init_uart

This function takes a pointer to a `uart_parameters` structure as an argument and initializes the relevant registers with the provided values. It is also responsible for powering on the peripheral through the register `PCONP`.

4.2 DMA Library

4.2.1 dma_transfer

```
1 struct dma_transfer{
2     uint32_t channel;
3     uint32_t size;
4     uint32_t dest;
5     uint32_t source;
6     uint32_t source_burst_size;
7     uint32_t dest_burst_size;
8     uint32_t source_increment;
9     uint32_t dest_increment;
10    uint32_t source_transfer_width;
11    uint32_t dest_transfer_width;
12    uint32_t terminal_count_enable;
13    uint32_t transfer_type;
14    uint32_t source_peripheral;
15    uint32_t dest_peripheral;
16    uint32_t lli_head;
17 };
```

Listing 4.2: DMA transfer struct

Just like with the UART definition, each value can either be a numeric value or a specific value expressed through a macro.

- **channel**: this value determines which channel we want to use. There are 8 channels available in the GPDMA, numbered from 0 to 7, with the lower numbers having higher priority. The macros `DMA_CHANNEL_x` can be used to fill this field.
- **size**: how many elements we want to transfer.
- **dest**: address of the destination.
- **source**: address of the source.
- **source_burst_size, dest_burst_size**: the burst size refers to the number of data items that are transferred in a single burst during a DMA transfer. It is typically measured in the number of items transferred per DMA transfer request. For example, a burst size of 8 means that 8 data items are transferred per DMA request. The burst size can affect the performance of the DMA transfer and may need to be optimized for the specific application. The values these two fields can assume are all the power of 2 between 1 and 256, using the macros `DMA_BURST_SIZE_x`.

- **source_increment, dest_increment:** these values determine if, after each transfer, we want to increment the source and destination address. The possible values are `DMA_NO_INCREMENT` and `DMA_INCREMENT`.
- **source_transfer_width, dest_transfer_width:** width of each element of source and destination. Possible values are 8 (byte), 16 (half word) and 32 (word), expressed through the macros `DMA_TRANSFER_WIDTH_8`, `DMA_TRANSFER_WIDTH_16` and `DMA_TRANSFER_WIDTH_32`.
- **terminal_count_enable:** with the value `DMA_TERMINAL_COUNT_ENABLE` we generate an interrupt at the end of the transaction, otherwise (using `DMA_TERMINAL_COUNT_NOT_ENABLE`) no.
- **transfer_type:** there are 4 kinds of transfer: memory to memory (`DMA_M2M`), memory to peripheral (`DMA_M2P`), peripheral to memory (`DMA_P2M`) and peripheral to peripheral (`DMA_P2P`).
- **source_peripheral, dest_peripheral:** these fields determine which are the source or destination peripheral of the transfer. This field is used to determine when to initiate a transaction based on the status of the peripheral. Given the large number of peripherals, we have not defined a macro for each one. Possible values can be found on page 592 of the user manual.
- **lli_head:** if we want to use DMA with the scatter gather technique, we fill this field with the pointer to the head of the linked list (whose structure is defined in `dma_lli_element`).

4.2.2 dma_lli_element

```

1 struct dma_lli_element{
2     uint32_t      src;
3     uint32_t      dest;
4     struct dma_lli_element*  next_lli;
5     uint32_t      ctrl;
6 };

```

Listing 4.3: DMA transfer struct

This structure is used to perform a scatter gather operation.

4.2.3 init_DMA

This function activates the GPDMA through the register `PCONP` and enables the associated interrupt.

4.2.4 transfer_DMA

This function uses a pointer to `dma_struct` as argument and initialize a DMA transaction.

4.3 Application

An application based on Micrium OS has been developed, in order to test the libraries and simulate the DMA attack on a real system, in this specific case the target board is the LandTiger V2.0 NXP LPC1768 ARM development board.

The simulated scenario of the application is a DMA attack to the control unit charged to monitor the state of the payload system in a spacecraft. The aim is to cause the failure of the mission, changing willingly the parameters of the sensors to shut down the on board computer OBC, which cannot operate over certain values of pressure and temperature.

4.3.1 Setup

The setup used for the development of the application has been the following:

- LandTiger V2.0 NXP LPC1768 ARM development board
- Keil uVision 5 running on a General Purpose PC
- Putty terminal emulator open source
- Uart to USB cable

4.3.2 Micrium OS

The application is built above the operating system Micrium, which is preemptive, deterministic, royalty-free and based on a flat architecture.

As already discussed in 2.2, DMA attacks are even easier to perform on these kind of architectures, because there is no distinct division between application and operating system, so it is possible to access reserved memory area, through the application.

Following the Task Model typical of Micrium, few tasks have been created and relative TCB and Stack have been allocated. These are:

- APP task
- Lock task
- Display task

4.3.3 APP task

The purpose of the application is to communicate with sensors and take actions depending on the data, so it is required to set-up UART communication and DMA transfer. This task has the highest priority.

In order to initialize a UART communication, a struct `uart_parameters` has to be defined and instantiated. Here all the parameters are set, in particular the most important ones are:

- **Clock frequency** : (CPU Clock)/4
- **Baud rate**: 9645 (Calculated using algorithm explained in user manual)
- **DMA mode**: ON
- **Trigger level**: 1 (as soon as an element arrives a DMA request is sent)
- **Word size**: 8 bits (1 Byte)

After the initialization of the UART connection the task is suspended and will be resumed by the `lock_task`.

As soon as `app_task` is resumed by `lock_task`, it initializes the DMA transfer, setting all the required parameters in an endless loop. The key ones are:

- **Destination**: Struct handling the data coming from the sensor
- **Source**: Receiver Buffer Register of UART
- **Burst size Dest/Source**: 1 (Just transfer one data item at a time)
- **Terminal Count Enable**: ON (Generate interrupt at the end of the transfer)
- **Transfer Type**: Peripheral to Memory (UART to memory)

After the parameters are set up, `transfer_DMA` function is called and the DMA transfer is initialized. Then the task is suspended again.

4.3.4 Lock task

Data coming from the sensor has a CRC field which is calculated before sending the data to the board. `Lock task` recalculates the CRC and compares it to the received one, to make sure that the data was sent correctly.

If the data is accepted, which means the calculated CRC corresponds to the received one, the condition to open the door is checked, updating the related flag. Depending on its final value the door is opened or closed and the state is saved.

The `App_task` is then resumed and the `Lock_task` is suspended.

4.3.5 Display task

This task has the lowest priority among the others, and it is responsible to display on screen the data coming from the sensor.

4.3.6 DMA_IRQHandler

This DMA handler is called when the DMA transfer is completed. The door's flag is reset and the LCD is cleared.

Depending on the result of the redundancy check, the data might be taken or not. If it is, no action is performed because `Lock_task` already did it, but if the CRC condition is not satisfied the door's state still needs to be decided and this is done restoring previous state.

Finally it resumes the `Lock_task`, restarting the flow of the application.

Bibliography

- [1] Markuze Alex, Shay Vargaftik, Gil Kupfer, Boris Pismeny, Nadav Amit, Adam Morrison, and Dan Tsafir. Characterizing, exploiting, and detecting dma code injection vulnerabilities in the presence of an iommu. In *Proceedings of the Sixteenth European Conference on Computer Systems*, EuroSys '21, page 395–409, New York, NY, USA, 2021. Association for Computing Machinery.
- [2] Ron Stajnsrod, Raz Yehuda, and Nezer Zaidenberg. Attacking trust-zone on devices lacking memory protection. *Journal of Computer Virology and Hacking Techniques*, 18, 09 2022.
- [3] Patrick Stewin and Iurii Bystrov. Understanding dma malware. In Ulrich Flegel, Evangelos Markatos, and William Robertson, editors, *Detection of Intrusions and Malware, and Vulnerability Assessment*, pages 21–41, Berlin, Heidelberg, 2013. Springer Berlin Heidelberg.
- [4] Dimitrios Tychalas, Anastasis Keliris, and Michail Maniatakos. Stealthy information leakage through peripheral exploitation in modern embedded systems. *IEEE Transactions on Device and Materials Reliability*, 20(2):308–318, 2020.